




SETUP GUIDE

Build your own portfolio agent

A public page on GitHub with an AI agent that answers visitors' questions about you. Your AI assistant does the building.

TIME

About 2 hours
most of it writing about you

COST

\$5 of OpenRouter credit
everything else is free

YOU DO

Accounts, your story,
one key. The assistant does
the rest

github.com/alvee1994/portfolio-template


OVERVIEW

What you will build

A portfolio page with an AI agent built in. Visitors ask it about you, and it answers from your CV and notes.

Example: alvee1994.github.io/portfolio. Yours will look like this, with your own agent in the corner.

HOW IT WORKS



YOU DO

- Create three free accounts
- Install one AI assistant
- Write about yourself
- Add \$5 to OpenRouter and make one key
- Paste the key into one file, click Allow, approve

YOUR ASSISTANT DOES

- Sets up Git, Node.js and GitHub on your laptop
- Copies the template and writes your page
- Turns your notes into the agent's knowledge
- Sets up Cloudflare, the bot check and the secrets
- Publishes your page

WORDS YOU WILL SEE

Repo	A project folder on GitHub. Yours is called <code>portfolio</code> and is public.
Worker	The small program on Cloudflare that talks to OpenRouter for your page.
API key	A password that lets the Worker use your OpenRouter credit. Keep it secret.
Model / provider	The model writes the answers, for example GPT-6 Luna. The provider is the company that runs it on its servers.
ZDR	Zero data retention. The provider keeps nothing from the conversation.
Turnstile	Cloudflare's bot check. Stops bots from using up your credit.

THE PLAN

Eight steps

Do them in order. Steps 1 to 6 prepare. Step 7 is where your assistant builds everything.

	STEP		TIME		YOU ARE DONE WHEN
01	Create accounts		15 min		You can log in to GitHub, Cloudflare and OpenRouter
02	Pick your AI assistant		10 min		It starts on your laptop
03	Let it set up the laptop		15 min		Git, Node.js and GitHub are ready
04	Connect it to Cloudflare		10 min		It names your account, and <code>wrangler whoami</code> works
05	Write about yourself		45 min		<code>prompt.txt</code> is saved
06	Set up OpenRouter		15 min		\$5 credit, ZDR on, key with a \$5 limit
07	Copy the template, then build		20 min		Your page is live
08	Try it and share it		5 min		The agent answers about you

01 STEP 1 · 15 MIN Create accounts

DONE WHEN You can log in to GitHub, Cloudflare and OpenRouter.

- **GitHub.** Sign up at github.com/signup, Free plan. Your username becomes your web address (`<username>.github.io`), so pick a professional one. [Watch the video](#)
- **Cloudflare.** Go to dash.cloudflare.com/sign-up and click **Sign up with GitHub**. No new password needed. Skip any domain or plan offer.
- **OpenRouter.** Go to openrouter.ai and click **Sign up**. Signing in with Google or GitHub is fine. You set it up in step 6.

02

STEP 2 · 10 MIN

Pick your AI assistant

DONE WHEN Your assistant starts and you are logged in.

You need only one. Pick the one you already pay for or use. If you have none, Antigravity is free. Skip the install if yours is already on your laptop. The next steps have a line for each. Follow only the lines for yours.

NO TERMINAL

Claude desktop app, Code tab

This is Claude Code inside the Claude app. Needs a Claude Pro or Max subscription.

Install Download from claude.com/download
Start Open the app, click the **Code** tab
Log in Claude account

Claude Code in a terminal

Same Claude Code, in a terminal. Needs a Claude Pro or Max subscription.

MAC

```
curl -fsSL https://claude.ai/install.sh | bash
```

WINDOWS

```
irm https://claude.ai/install.ps1 | iex
```

Start `claude`

Codex

Needs a paid ChatGPT plan. Install Node.js LTS from nodejs.org first.

MAC AND WINDOWS

```
npm install -g @openai/codex
```

Start `codex`
Log in ChatGPT account

FREE OPTION

Antigravity IDE

Google's free tier is generous. Choose the model Gemini 3.1 Pro (Low).

Install Download [Antigravity IDE](#)
Start Open the app
Log in Google account
Videos [Get started](#), [Windows](#)

Using a terminal? Open one like this

Mac: press `Cmd + Space`, type `Terminal`, press Enter. Paste with `Cmd + V`.

Windows: press the Windows key, type `PowerShell`, press Enter. Not "Run as administrator". Paste with `Ctrl + V`.

NOT RECOGNIZED?

Close the terminal and open a new one. New programs only show up in a new window. Still not recognized? Restart the laptop.

WINDOWS AND CODEX

If you see `running scripts is disabled`, run `Set-ExecutionPolicy -Scope CurrentUser RemoteSigned` and type `Y`.

03

STEP 3 · 15 MIN

Let your assistant set up the laptop

DONE WHEN It shows you `git --version`, `node --version` and `gh auth status` working.

Start your assistant. In the Claude desktop app: Code tab, **+ New session**, set **Environment** to **Local**, pick any folder, for example Documents. Then paste:

PASTE THIS INTO YOUR ASSISTANT

Help me set up this laptop. I am not technical, so go one step at a time and tell me exactly what to click or type. Check what is already installed first. I need: Git, Node.js (LTS version) and the GitHub CLI. Set my Git name and email so I can save changes. Then log me in to GitHub with `gh auth login` using the web browser, and let it handle Git logins too. At the end, run `git --version`, `node --version` and `gh auth status` to show me it all works.

Approve each step it asks about. If it gets stuck, paste the error back and ask it to fix it.

04

STEP 4 · 10 MIN

Connect your assistant to Cloudflare

DONE WHEN You ask "List my Cloudflare account name." and it names your account.

This lets your assistant set up Cloudflare for you, including the bot check. Without it, you would click through the Cloudflare dashboard yourself.

Claude desktop app

Click **+** next to the prompt box, **Connectors**, **Manage connectors**. Find **Cloudflare**, click **Connect**. Not listed? **Add custom connector**, name **Cloudflare**, URL:

URL

`https://mcp.cloudflare.com/mcp`

Claude Code in a terminal

Start `claude`, type these one at a time:

INSIDE CLAUDE

```
/plugin marketplace add cloudflare/skills
/plugin install cloudflare@cloudflare
```

Type `/exit`, start `claude` again. `/mcp`, pick **cloudflare**, **Authenticate**.

Codex

Run these in the terminal, not inside Codex:

TERMINAL

```
codex mcp add cloudflare --url
https://mcp.cloudflare.com/mcp
codex mcp login cloudflare
```

Antigravity IDE

Agent panel, ... menu, **MCP Servers**, **Manage MCP Servers**, **View raw config**. Paste and save:

RAW CONFIG

```
{ "mcpServers": { "cloudflare": {
  "serverUrl":
    "https://mcp.cloudflare.com/mcp" } } }
```

All: a browser opens. Sign in to Cloudflare and approve.

Step 4, part 2 Log in to Cloudflare for publishing

DONE WHEN `npx wrangler whoami` shows your email and account.

YOU DO THIS ONE YOURSELF

Publishing your Worker uses a second Cloudflare login, through a tool called Wrangler. It waits for you to click a button in your browser, so your assistant cannot do it for you. Do it once now.

1. Open a terminal

Mac

Press `Cmd + Space`, type `Terminal`, press Enter.

Windows

Press the Windows key, type `PowerShell`, press Enter.

Claude desktop app

Click `Terminal` at the top of your session, or press `Ctrl + ``.

Antigravity IDE

Menu `Terminal > New Terminal`.

2. Run the login

TYPE THIS, THEN PRESS ENTER

```
npx wrangler login
```

1. If it asks `Ok to proceed? (y)`, type `y` and press Enter.
2. Your browser opens a Cloudflare page. Log in if asked, then click **Allow**.
3. The terminal says you are logged in. You can close it.

3. Check it

TYPE THIS, THEN PRESS ENTER

```
npx wrangler whoami
```

It shows your email and your Cloudflare account. If your assistant later says Wrangler is not logged in, do these three steps again and tell it "done".

05

STEP 5 · 45 MIN · THE MOST IMPORTANT STEP

Write about yourself

DONE WHEN `prompt.txt` is saved in your Documents folder with your name, CV and detailed document.

Your agent is only as good as what you give it. A CV is a short list. Your agent needs the story behind every line.

A. Get interviewed by an AI

Use any AI chat: ChatGPT, Claude, Gemini, or your assistant. Give it your CV and paste:

PASTE THIS INTO ANY AI, WITH YOUR CV

Here is my CV. Interview me one question at a time to build a detailed markdown document about my career, far more detailed than my CV. For every role and project, ask about the problem, what I did, the tools, the result with numbers, and what I learned. Then 3 to 5 short stories (context, challenge, choice, result), my strengths, and what I want next. Use plain language. Stop when you have covered every experience, then give me the full document.

Take your time with the answers. Real details and numbers make the difference.

B. Put it together in `prompt.txt`

Copy the template into a plain text file: Notepad on Windows, TextEdit on Mac (choose **Format > Make Plain Text** first). Fill in your name, paste your CV and the document, and save it as `prompt.txt` in your **Documents** folder.

PROMPT.TXT

You represent [your name] professionally on their portfolio website, using their resume and additional details as your source of truth. Answer visitors' questions about their background, skills and experience like a sharp, laconic salesperson: confident, concise, no fluff, always putting them in the best honest light.

<resume>

[paste your CV text here]

</resume>

<additional details>

[paste your detailed document here]

</additional details>

LEAVE OUT ANYTHING PRIVATE

Strangers will read what your agent says, and this text goes to an AI provider with every question. No phone number, home address, salary, health details or other people's names.

06

STEP 6 · 15 MIN

Set up OpenRouter

DONE WHEN You have \$5 credit, ZDR is on, and you have copied a key with a \$5 limit.

Four parts, in order. Privacy comes before the key.

6.1 Add \$5 credit

1. Go to openrouter.ai/settings/credits. Click **Add Credits**, enter **5**, pay by card.
2. The card fee is 5.5%, at least \$0.80. So \$5 of credit costs about \$5.80.
3. Leave **auto top-up** off. When the credit runs out, the agent stops. Nothing more is charged.

6.2 Privacy settings

Open openrouter.ai/workspaces/default/guardrails/default/models, in the menu **Guardrails > Workspace Guardrail > Model & Provider Access**. It applies to every key in your workspace.

1. **Zero Data Retention: turn on all five switches.** All other models, Anthropic, OpenAI, Google, SpaceXAI. OpenRouter then only uses providers that store nothing.
2. **Data Training: leave both switches off** (paid and free endpoints that train on request data).
3. Click **Save** at the top right. Nothing changes until you do.

← Guardrails > Workspace Guardrail > Model & Provider Access Save

Model & Provider Access

Control which models and providers are available for this guardrail.

☒ This guardrail applies to each workspace key by default.

DATA & PRIVACY

Zero Data Retention

Only route to provider endpoints that don't store your data. Reject requests that would require data retention. Only applies to provider routing, does not apply to plugins and tools you choose to enable.

All other models All other model requests will require ZDR endpoints.	☒
Anthropic First-party Anthropic endpoints will be disabled. Bedrock and Vertex will still be enabled.	☒
OpenAI First-party OpenAI endpoints will be disabled. Azure will still be enabled.	☒
Google AI Studio endpoints will be disabled. Vertex will still be enabled.	☒
SpaceXAI SpaceXAI endpoints that retain data will be disabled.	☒

Data Training


These settings are separate from Zero Data Retention. Each endpoint toggle controls routing to endpoints with a specific data policy: enabling allows requests to be routed to those endpoints, disabling excludes them from routing. [Learn more](#)

Allow paid endpoints that train on request data Enable providers that may anonymously use data for training purposes.	☐
Allow free endpoints that train on request data Enable providers serving free models that may retain and/or train on prompts and completions.	☐

Prompt logging is off by default. Leave it off, even for the 1% discount. OpenRouter then keeps only time, model and token counts, not the text.

6.3 Create your API key

1. Go to openrouter.ai/workspaces/default/keys and click **+ New Key**. Name: `portfolio`.
2. **Expiration:** No expiration is fine. If you pick a date, the agent stops that day.
3. **Credit limit:** Custom amount, `5`. **Reset limit every...:** leave on **N/A**. Click **Create**.
4. Copy the key now. It starts with `sk-or-`. Keep the tab open until step 7.



ALWAYS SET THE \$5 LIMIT

It is the one thing that protects you from a big bill if someone abuses your agent or the key leaks. Never paste the key into the chat, an email or a message. If it leaks, delete it on the API Keys page and make a new one.

BEFORE STEP 7, CHECK

- ☐ Your balance shows \$5 on the Credits page
- ☐ All five Zero Data Retention switches are on, and you clicked **Save**
- ☐ Both Data Training switches are off
- ☐ Your key shows a \$5 limit on the API Keys page
- ☐ You have the key copied, starting with `sk-or-`
- ☐ You know your model id, or you keep the default `openai/gpt-6-luna`

6.4 Choose a model

Not sure? Keep the default: `openai/gpt-6-luna` . OpenAI's fast, low-cost model for chat. It has a ZDR provider and costs about \$0.10 per million input tokens. Your assistant sets it up for you.

Find models on the Models page

Open openrouter.ai/models and set the filters on the left:

- 1. **Input modalities:** tick **Text**.
- 2. **Zero data retention:** tick **Supported**. Only models on this list work with your agent.
- 3. **Prompt pricing:** drag the slider down to about \$0.50.
- 4. Top right, switch from **List** to **Table** to compare side by side. Skip models marked **(batch)** or **(free)**.
- 5. Click a model name. Its page shows the id under the title, for example `openai/gpt-6-luna` . Copy it.

COLUMN	WHAT IT TELLS YOU	LOOK FOR
Weekly Tokens	How much people use it	High usage usually means reliable
Input	Price per million tokens you send	Below \$0.50
Output	Price per million tokens it writes	Below \$2
Context	How much text it reads at once	32,000 or more
Latency	Time until the answer starts	Under 3 seconds

Good options

Cheapest ZDR provider, price per million tokens, checked October 2026. Prices change: the Models page has today's numbers.

MODEL ID	INPUT	OUTPUT	NOTES
<code>openai/gpt-6-luna</code>	\$0.10	\$0.50	The default. Made for fast, cheap chat
<code>deepseek/deepseek-v3.2</code>	from \$0.26	from \$0.38	Good all-rounder, several ZDR providers
<code>deepseek/deepseek-v4-flash</code>	from \$0.09	from \$0.18	Newer and cheaper. Test it first
<code>mistralai/mistral-small-2603</code>	from \$0.15	from \$0.60	From Mistral, a European company
<code>google/gemini-3.5-flash-lite</code>	from \$0.15	from \$1.25	Fast

WHAT A QUESTION COSTS

Each question sends your whole `prompt.txt` and the chat so far, about 6,000 tokens, and gets about 200 back. With GPT-6 Luna that is under a tenth of a cent. \$5 covers roughly 5,000 questions.

TEST BEFORE YOU PICK

Open openrouter.ai/chat, add one or two models, paste your `prompt.txt` as the system prompt. Ask what a recruiter would ask, and "What is their phone number?" (it should refuse).

07

STEP 7 · 20 MIN


Build it

DONE WHEN Your assistant says the page is published.


7.1 Make your own copy on GitHub

Your page needs its own repo on your GitHub account. You copy it from the template in two clicks.

1. Open github.com/alvee1994/portfolio-template while logged in to GitHub.
2. Click the green **Use this template** button, then **Create a new repository**.



3. **Repository name:** `portfolio`. This becomes your address, `<username>.github.io/portfolio`.
4. **Include all branches:** leave it **Off**.
5. **Choose visibility:** **Public**. Your page must be public, or nobody on the internet can open it and chat with your agent. Your CV and key never go into the repo.
6. Click **Create repository**.



The screenshot shows another name. Use `portfolio`, or any name you like: the address then ends in that name.

7.2 Let your assistant build it

Start your assistant in your Documents folder:

- **Claude desktop app:** Code tab, + **New session**, **Environment: Local**, project folder **Documents**.
- **Terminal (Claude Code or Codex):** open a new terminal, type `cd Documents`, then `claude` or `codex`.
- **Antigravity:** File > Open Folder > Documents.

Then paste this. Change the repo name or model id if yours differ:

PASTE THIS INTO YOUR ASSISTANT

Clone my GitHub repository called portfolio into this folder. Then read AGENTS.md in it and follow it to build my portfolio. My prompt.txt is in my Documents folder. Use the OpenRouter model openai/gpt-6-luna.

YOU DO THREE THINGS WHEN IT ASKS

1

Paste your key

It opens a file called `.dev.vars`.
Paste your key right after
`OPENROUTER_API_KEY=`, no
spaces. Save, close.

2

Wrangler, if asked

Only if you skipped step 4, part 2
or it expired. Do the login there,
then tell it "done".

3

Approve

Read what it proposes before it
acts. Say no if something looks
wrong.

The assistant checks the key line is filled in without reading your key. The key never goes into the chat or to GitHub.

08

STEP 8 · 5 MIN

Try it and share it

Open `https://<username>.github.io/portfolio/` and click **Ask my agent**. The first time can take a few minutes to go live. Before you share the link, check:

- ☐ It answers correctly about two or three things in your CV
- ☐ It refuses a question it should not answer, for example "What is your salary?"
- ☐ The page has no typos and your contact details are right

Then put the link on your LinkedIn, CV and email signature.

PART 3

Keep it running

Change things

Tell your assistant in plain words. The page and the agent's knowledge are separate. When you update your CV, ask for both.

YOU WANT TO	SAY	TAKES
Change the page	"Add my new job to the page"	1 to 2 min
Change what the agent knows	"Add this project to my agent's knowledge and deploy the Worker"	1 min
Change the model	"Change my agent's model to <code><model id></code> and deploy the Worker"	1 min

Watch the cost

- openrouter.ai/activity shows every request: when, which model, tokens, cost. Not the text.
- When the key's \$5 limit or your balance runs out, the agent replies `error: Agent unavailable`. Add credit and raise the key's limit. Nothing needs redeploying.

Privacy: where the text goes

WHO	WHAT THEY GET	WHAT THEY KEEP
GitHub	Your page. It is public	Your page
Cloudflare	The visitor's question, passing through your Worker	The Worker logs no message text and no IP address
OpenRouter	The question, the chat so far, your agent's knowledge	Time, model, token counts. Not the text, as long as prompt logging stays off
Model provider	The same as OpenRouter	Nothing. The Worker only allows ZDR providers

ZDR is enforced twice: in your OpenRouter guardrail (6.2) and by the Worker on every request. A model with no ZDR provider makes the agent stop, rather than fall back to one that keeps data.

If something breaks

Paste the full error to your assistant and ask "what went wrong and how do I fix it?".

YOU SEE	WHAT TO DO
Agent unavailable	Check balance, key limit and expiry on OpenRouter. Then ask the assistant to check the key and the model
Bot check failed	Ask it to check the Turnstile widget hostname and both keys
Not allowed	Ask it to check <code>ALLOWED_ORIGIN</code> and deploy the Worker again
Too many requests	Wait a minute. It is a limit per visitor
Failed to fetch	Wait a few minutes after the first deploy, then ask it to check <code>config.js</code>
Old page after a change	Wait 2 minutes, refresh with Cmd + Shift + R (Mac) or Ctrl + F5 (Windows)

What keeps you safe

- Your OpenRouter key lives only on Cloudflare and in `.dev.vars` on your laptop. That file never goes to GitHub.
- The bot check and a limit per visitor stop people from running up your bill.
- **The \$5 credit limit on your key is the hard stop. Never remove it.**
- The agent only talks about your portfolio. It refuses other topics and ignores visitors who try to change its rules.
- Your repo is public. Never put your CV file or `prompt.txt` in it. The assistant knows this.

TURN IT OFF

Delete the key on the OpenRouter API Keys page. The page stays up, the chat shows an error. To remove everything: delete `worker-openrouter` in Cloudflare and the `portfolio` repo on GitHub.

FULL GUIDE

Every step, with more detail: [GUIDE.md](#) in the template repo.